



Joel on Software

The Joel Test: 12 Steps to Code

by Joel Spolsky

Wednesday, August 09, 2000

Have you ever heard of [SEMA](#)? It's a fairly esoteric s

measuring how good a software team is. No, *wait! D link!* It will take you about six years just to *understa*. I've come up with my own, highly irresponsible, slop quality of a software team. The great part about it is 3 minutes. With all the time you save, you can go to :

The Joel Test

1. Do you use source control?
2. Can you make a build in one step?
3. Do you make daily builds?
4. Do you have a bug database?
5. Do you fix bugs before writing new code?
6. Do you have an up-to-date schedule?
7. Do you have a spec?
8. Do programmers have quiet working conditions?
9. Do you use the best tools money can buy?
10. Do you have testers?
11. Do new candidates write code during the interview?
12. Do you do hallway usability testing?

The neat thing about The Joel Test is that it's easy to answer **no** to each question. You don't have to figure out how many bugs per day or average-bugs-per-inflection-point. Give your

each "yes" answer. The bummer about The Joel Test *shouldn't* use it to make sure that your nuclear power is safe.

A score of 12 is perfect, 11 is tolerable, but 10 or lower indicates serious problems. The truth is that most software organizations are running with a score of 2 or 3, and they need *serious* changes. Companies like Microsoft run at 12 full-time.

Of course, these are not the only factors that determine success or failure: in particular, if you have a great software team that produces a product that nobody wants, well, people aren't going to buy it. It's possible to imagine a team of "gunslingers" that does all of this stuff that still manages to produce incredible software that changes the world. But, all else being equal, if you get the Joel Test right, you'll have a disciplined team that can consistently

1. Do you use source control?

I've used commercial source control packages, and I've used CVS, which is free, and let me tell you, CVS is *fine*. But if you don't use source control, you're going to stress out trying to get everyone to work together. Programmers have no way to know what you did. Mistakes can't be rolled back easily. The other nice thing about source control systems is that the source code itself is not on every programmer's hard drive -- I've never heard of a source control system that lost a lot of code.

2. Can you make a build in one step?

By this I mean: how many steps does it take to make

from the latest source snapshot? On good teams, the you can run that does a full checkout from scratch, r of code, makes the EXEs, in all their various versions #ifdef combinations, creates the installation package final media -- CDROM layout, download website, wh

If the process takes any more than one step, it is pro when you get closer to shipping, you want to have a fixing the "last" bug, making the final EXEs, etc. If it compile the code, run the installation builder, etc., y crazy and you're going to make silly mistakes.

For this very reason, the last company I worked at sw WISE to InstallShield: we *required* that the installat able to run, from a script, automatically, overnight, u scheduler, and WISE couldn't run from the schedule threw it out. (The kind folks at WISE assure me that version does support nightly builds.)

3. Do you make daily builds?

When you're using source control, sometimes one pr accidentally checks in something that breaks the bui they've added a new source file, and everything comp machine, but they forgot to add the source file to the So they lock their machine and go home, oblivious a nobody else can work, so they have to go home too, u

Breaking the build is so bad (and so common) that it daily builds, to insure that no breakage goes unnotic

teams, one good way to insure that breakages are fixed is to do the daily build every afternoon at, say, lunchtime. Do many checkins as possible before lunch. When the daily build is done. If it worked, great! Everybody checks out a new version of the source and goes on working. If the build fails, it's not a problem, but everybody can keep on working with the pre-lunch version of the source.

On the Excel team we had a rule that whoever broke the build, "punishment", was responsible for babysitting the build until someone else broke it. This was a good incentive not to break the build, and a good way to rotate everyone through the job so that everyone learned how it worked.

Read more about daily builds in my article [Daily Builds for the Friend](#).

4. Do you have a bug database?

I don't care what you say. If you are developing code without one, without an organized database listing all known bugs, you are going to ship low quality code. Lots of people think they can hold the bug list in their heads. Nonsequence. You can't remember more than two or three bugs at a time, and by the morning, or in the rush of shipping, they are forgotten. You have to keep track of bugs formally.

Bug databases can be complicated or simple. A minimal bug database must include the following data for every bug:

- complete steps to reproduce the bug
- expected behavior
- observed (buggy) behavior
- who it's assigned to
- whether it has been fixed or not

If the complexity of bug tracking software is the only thing that keeps you from tracking your bugs, just make a simple 5 column spreadsheet with these crucial fields and *start using it*.

For more on bug tracking, read [Painless Bug Tracking](#)

5. Do you fix bugs before writing new code

The very first version of Microsoft Word for Windows was a "death march" project. It took forever. It kept slipping. The team was working ridiculous hours, the project was cancelled again, and again, and the stress was incredible. When it finally shipped, years late, Microsoft sent the whole team to Cancun for a vacation, then sat down for some serious

What they realized was that the project managers had been insistent on keeping to the "schedule" that programmers had rushed through the coding process, writing extremely buggy code because the bug fixing phase was not a part of the formal process. There was no attempt to keep the bug-count down. (The story goes that one programmer, who had to write a routine to calculate the height of a line of text, simply wrote "return 1" and waited for the bug report to come in about how his function was always correct. The schedule was merely a checklist of

to be turned into bugs. In the post-mortem, this was "infinite defects methodology".

To correct the problem, Microsoft universally adopted what was called a "zero defects methodology". Many of the people at the company giggled, since it sounded like management could reduce the bug count by executive fiat. Actually, it meant that at any given time, the highest priority is to fix bugs *before* writing any new code. Here's why.

In general, the longer you wait before fixing a bug, the more time (and money) it is to fix.

For example, when you make a typo or syntax error and it gets caught, fixing it is basically trivial.

When you have a bug in your code that you see the first time you run it, you will be able to fix it in no time at all, because it's still fresh in your mind.

If you find a bug in some code that you wrote a few days ago, it might take you a while to hunt it down, but when you reread the code you wrote, you'll remember everything and you'll be able to fix it in a reasonable amount of time.

But if you find a bug in code that you wrote a few months ago, you probably have forgotten a lot of things about that code, and it's harder to fix. By that time you may be fixing somebody else's code, and they may be in Aruba on vacation, in which case

And if you find a bug in code that has *already shipped*,
incur incredible expense getting it fixed.

What this means is that if you have a schedule with a lot of bugs remaining to be fixed, the schedule is unreliable. But if you fix the *known* bugs, and all that's left is new code, then you can be stunningly more accurate.

Page 8 of 17

6. Do you have an up-to-date schedule?

Which brings us to schedules. If your code is at all in business, there are lots of reasons why it's important to know when the code is going to be done. Programmers are crabby about making schedules. "It will be done when we scream at the business people."

Unfortunately, that just doesn't cut it. There are too many decisions that the business needs to make well in advance of the code: demos, trade shows, advertising, etc. And that's why this is to have a schedule, and to keep it up to date.

The other crucial thing about having a schedule is to decide what features you are going to do, and then to cut the least important features and *cut them* rather than [featuritis](#) (a.k.a. scope creep).

Keeping schedules does not have to be hard. Read my [Software Schedules](#), which describes a simple way to keep schedules.

7. Do you have a spec?

Writing specs is like flossing: everybody agrees that it's important, but nobody does it.

I'm not sure why this is, but it's probably because most programmers hate writing documents. As a result, when teams confront a problem, programmers attack a problem, they prefer to express their solution in code, rather than in documents. They would much rather

write code than produce a spec first.

At the design stage, when you discover problems, you fix them easily by editing a few lines of text. Once the code is written, the cost of fixing problems is dramatically higher, both emotionally (hate to throw away code) and in terms of time, so they don't actually fix the problems. Software that wasn't built with a spec usually winds up badly designed and the schedule gets out of control. This seems to have been the problem at Netscape, where several versions grew into such a mess that management [suggested](#) to throw out the code and start over. And then they made it over again with Mozilla, creating a monster that spun out of control and took *several years* to get to alpha stage.

My pet theory is that this problem can be fixed by teaching programmers to be less reluctant writers by sending them [an intensive course in writing](#). Another solution is to hire program managers who produce the written spec. In any case, we should enforce the simple rule "no code without spec".

Learn all about writing specs by reading my [4-part series](#).

8. Do programmers have quiet working conditions?

There are extensively documented productivity gains from giving knowledge workers space, quiet, and privacy. The software management book [Peopleware](#) documents these benefits extensively.

Here's the trouble. We all know that knowledge workers

getting into "flow", also known as being "in the zone" fully concentrated on their work and fully tuned out environment. They lose track of time and produce great work in absolute concentration. This is when they get all of their work done. Writers, programmers, scientists, and even athletes will tell you about being in the zone.

The trouble is, getting into "the zone" is not easy. When we measure it, it looks like it takes an average of 15 minutes to get working at maximum productivity. Sometimes, if you've already done a lot of creative work that day, you just don't get into the zone and you spend the rest of your work day fiddling with the web, playing Tetris.

The other trouble is that it's so easy to get knocked out of the zone. Noise, phone calls, going out for lunch, having to drive to Starbucks for coffee, and interruptions by coworkers -- all knock you out of the zone. Even a 1 minute interruption by a coworker asking you a question, causing a 1 minute interruption, this knocks you out of the zone badly enough that it takes an hour to get productive again, your overall productivity is a lot of trouble. If you're in a noisy bullpen environment like the ones caffeinated dotcoms love to create, with marketing guys and the phone next to programmers, your productivity will be a lot of trouble. Knowledge workers get interrupted time after time and can't get into the zone.

With programmers, it's especially hard. Productivity is a lot of trouble. Being able to juggle a lot of little details in short term memory is a lot of trouble.

Any kind of interruption can cause these details to come down. When you resume work, you can't remember (like local variable names you were using, or where you were implementing that search algorithm) and you have to pick these things up, which slows you down a lot until you get back to speed.

Here's the simple algebra. Let's say (as the evidence suggests) that if we interrupt a programmer, even for a minute, it's blowing away 15 minutes of productivity. For this example, let's take two programmers, Jeff and Mutt, in open cubicles next to each other in a standard Dilbert veal-fattening farm. Mutt can't remember the name of the Unicode version of the strcpy function. He has to look it up, which takes 30 seconds, or he could ask Jeff, who knows the answer in 15 seconds. Since he's sitting right next to Jeff, he asks Jeff, who is not distracted and loses 15 minutes of productivity (to save 30 seconds).

Now let's move them into separate offices with walls. When Mutt can't remember the name of that function, he has to look it up, which still takes 30 seconds, or he could ask Jeff, who is 45 feet away. It takes 45 seconds and involves standing up (not an easy task for the average physical fitness of programmers!). So he looks at Jeff, who is 45 feet away. Mutt loses 30 seconds of productivity, but we save 15 seconds. Ahhh!

9. Do you use the best tools money can buy?

Writing code in a compiled language is one of the last things that can't be done instantly on a garden variety home computer.

compilation process takes more than a few seconds, and greatest computer is going to save you time. If c even 15 seconds, programmers will get bored while t and switch over to reading [The Onion](#), which will suc hours of productivity.

Debugging GUI code with a single monitor system is impossible. If you're writing GUI code, two monitors much easier.

Most programmers eventually have to manipulate bi toolbars, and most programmers don't have a good k available. Trying to use Microsoft Paint to manipulat joke, but that's what most programmers have to do.

At [my last job](#), the system administrator kept sendin spam complaining that I was using more than ... get megabytes of hard drive space on the server. I pointe the price of hard drives these days, the cost of this sp significantly less than the cost of the *toilet paper* I us even 10 minutes cleaning up my directory would be : of productivity.

Top notch development teams don't torture t programmers. Even minor frustrations caused by underpowered tools add up, making programmers g unhappy. And a grumpy programmer is an unprodu

To add to all this... programmers are easily bribed by

coolest, latest stuff. This is a far cheaper way to get the latest stuff than actually paying competitive salaries!

10. Do you have testers?

If your team doesn't have dedicated testers, at least one or two programmers, you are either shipping buggy code or three programmers, you are either shipping buggy code or you're wasting money by having \$100/hour programmers that can be done by \$30/hour testers. Skimping on testing is an outrageous false economy that I'm simply blown away by. People don't recognize it.

Read [Top Five \(Wrong\) Reasons You Don't Have Testers](#) I wrote about this subject.

11. Do new candidates write code during the interview?

Would you hire a magician without asking them to show you magic tricks? Of course not.

Would you hire a caterer for your wedding without tasting the food? I doubt it. (Unless it's Aunt Marge, and she would have you didn't let her make her "famous" chopped liver casserole.)

Yet, every day, programmers are hired on the basis of a good résumé or because the interviewer enjoyed chatting with them. They are asked trivia questions ("what's the difference between `CreateDialog()` and `DialogBox()`?) which could be answered by looking at the documentation. You don't care if they know thousands of trivia about programming, you care if they can write code.

produce code. Or, even worse, they are asked "AHA!" kind of questions that seem easy when you know the answer, but if you don't know the answer, they are impossible.

Please, just *stop doing this*. Do whatever you want do but make the candidate *write some code*. (For more [Guerrilla Guide to Interviewing](#).)

12. Do you do hallway usability testing?

A *hallway usability test* is where you grab the next person by in the hallway and force them to try to use the code. If you do this to five people, you will learn 95% of what you will learn about usability problems in your code.

Good user interface design is not as hard as you would think. It is crucial if you want customers to love and buy your product. Please read my [free online book on UI design](#), a short primer for programmers.

But the most important thing about user interfaces is to show your program to a handful of people, (in fact, five or six) and they will quickly discover the biggest problems people are having. See [Jakob Nielsen's article](#) explaining why. Even if your UI is lacking, as long as you force yourself to do hallway usability testing, which cost nothing, your UI will be much, much better.

Next: [Wasting Money on Cats](#)

Want to know more? You're reading [Joel on Software](#), years and years of completely raving mad articles about development, managing software teams, designing and running successful software companies, and rubber ducking.

About the author. I'm [Joel Spolsky](#), co-founder of [Fog Creek Software](#), a New York company that proves that you can pay programmers well and still be highly profitable. Programmers have private offices, free lunch, and work 40 hours a week. We pay for software if they're delighted. We make FogBugz, an enlightened [bug tracker](#) designed to help great teams make great software, Kiln, which simplifies source control and [cvs](#), and Fog Creek Copilot, which makes [remote desktop control](#). I'm also the co-founder of [Stack Overflow](#).

© 2000-2011 Joel Spolsky

joel@joelonsoftware.com

